

Learning-Enabled Verification of Distributed Systems with End-to-End Proofs

Federico Mora, University of California, Berkeley

Available at federico.morarocha.ca/quals. Navigate using  keys

(use Firefox and adjust size by zooming)

Distributed System Bugs are Common and Costly

Distributed bugs... are often severe... Not only are these outages widespread and expensive, they can be caused by bugs that were deployed to production months earlier... [For example] [amazon.com] went down because one remote server couldn't display any product information.

- Jacob Gabrielson, Senior Principal Engineer at Amazon Web Services

Proposal: Verification Methodology for Async Distributed Systems

Our collection of tools and techniques will be used to

- Verify message-passing distributed systems
 - for any number of nodes and any execution length
 - starting with asynchronous
- Without abstracting away messages
 - staying close to implementations (inside an existing programming language!)
- While giving meaningful feedback
 - generating end-to-end proofs on success
 - suggesting lemmas on failure
- With a high degree of automation
 - combining automated theorem proving and inductive reasoning

Proposal: Verification Methodology for Async Distributed Systems

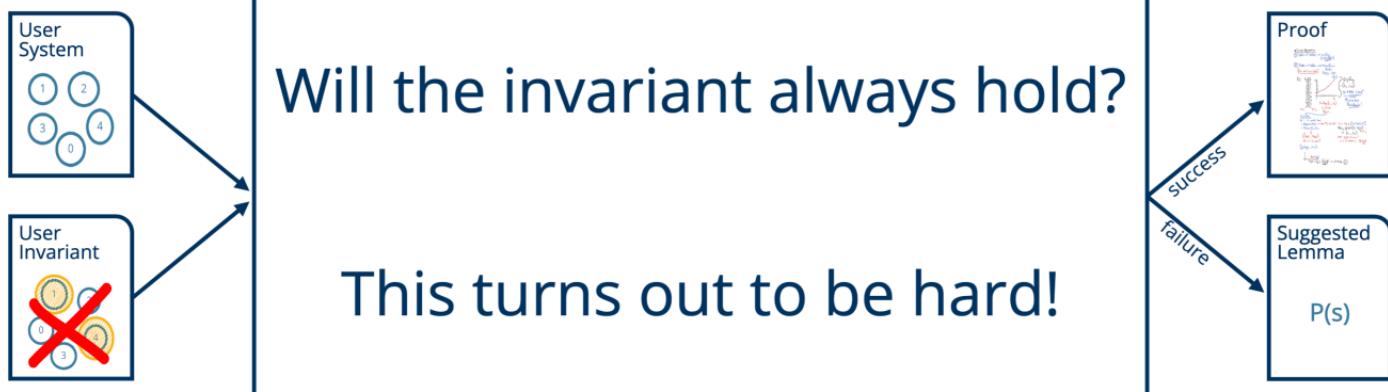


Table of Desired Features

	Domain-Specific	Implementation Level	Modular Verification	High Automation	Proof Generation	Specification Mining
TLA+	✓	✗	✓	✗	✗	✗
Ivy	✓	✗	✓	✓	✗	✗
mypyvvy	✓	✗	✗	✓	✗	✗
Coq/Lean/...	✗	✓	✓	✗	✓	✗
Dafny/F*/...	✗	✓	✓	✓	✗	✗
Proposal	✓	✓	✓*	✓	✓*	✓

Background

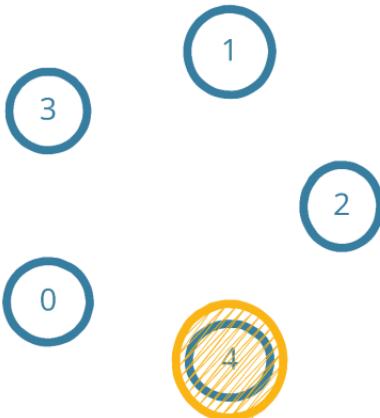
Asynchronous Systems and Verification

Reasoning About Messages

Coming Up with True Invariants

Future work: Extensions, Proofs, and Cloud Solving

Async System Example



Ring Leader Election Protocol

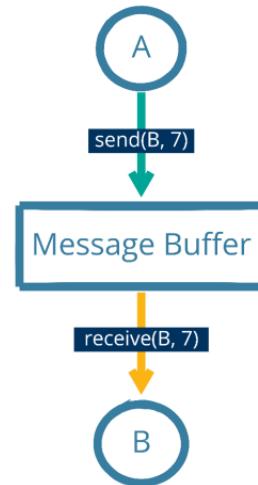
- Send your own identifier clockwise
- If you receive a value greater than your own, send new value clockwise
- If you receive your own identifier, declare yourself the winner

Log `send(to, payload) | receive(from, payload) | win(node)`

`send(4, 2); send(0, 4); send(0, 4); receive(0, 4) and send(3, 4);
send(3, 0); receive(0, 4) and send(3, 4); receive(3, 4) and send(1,
4); send(1, 3); receive(1, 4) and send(2, 4); receive(3, 4) and
send(1, 4); receive(2, 4) and send(4, 4); receive(4, 4) and win(4).`

Async Systems, Generally

- Nodes are state machines with a unique identifier
- The set of node identifiers N has at least two elements
- Nodes perform atomic actions
- Messages are pairs (t, p) , where
 - $t \in N$ represents the message target and
 - p is the message payload
- There is multiset of messages, called the message buffer
- Nodes can only read/write to their own state
- Except that every node n can (for any t and p)
 - $\text{send}(t, p)$ in the message buffer, or
 - $\text{receive}(n, p)$ from the message buffer.
 - If (n, p) is present, then $\text{receive}(n, p)$ returns p
 - Otherwise, $\text{receive}(n, p)$ returns \emptyset



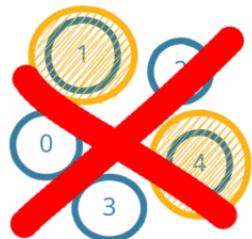
Verification Goal Example

Ring System Assumptions:

- The set of node identifiers is an unknown subset of the natural numbers
- Each node knows the identifier of the next "clockwise" node
- Initially, no node is in the winning state and the message buffer is empty

Ring System Specification (Verification Goal):

- There is always at most one node in the winning state



Verification Goal, Generally

- An invariant is a predicate over the state of the system
- We want to guarantee that an invariant will hold over every state of every possible execution
- The state of the system is the
 - state of every node, and
 - the state of the message buffer
- A system step is a pair of states (σ, σ') , where σ' is the result of a node's action on σ
- An execution is a sequence of states $\sigma_1 \dots \sigma_n$ such that
 - σ_1 is an initial state, and
 - for each i , (σ_i, σ_{i+1}) is a step of the system

Verification Technique: Proof by Induction

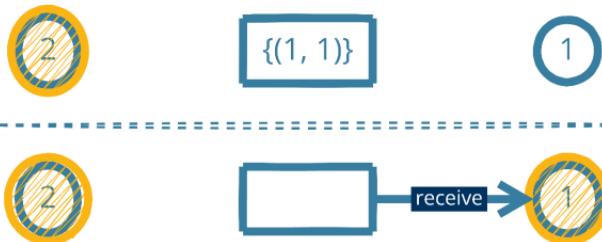
- An invariant I holds for every execution if
 1. Base case: $I(\sigma)$ holds for every initial state σ
 2. Inductive step: $I(\sigma)$ implies $I(\sigma')$ holds for every step (σ, σ')
- An invariant that meets condition (2) is inductive
- A counterexample to (2) is a counterexample to induction

Verification Challenges Example

Ring System Specification (Verification Goal):

- There is always at most one node in the winning state

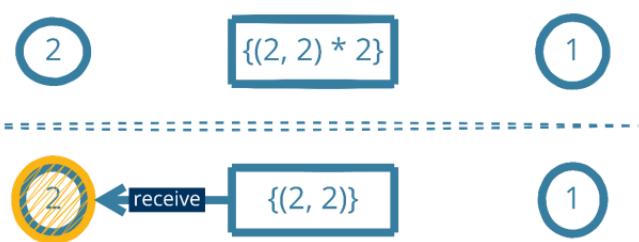
Counterexample to Induction (Failure):



Ring System Specification (Take 2):

- There is always at most one node in the winning state, and
- if a node won, then there is no message in the buffer whose target and payload are equal

Counterexample to Induction (Take 2):



Verification Challenges, Generally

Reasoning about "stray" messages.

- Some move to a more abstract model, e.g., mypyv's shared memory
- Some perform painstaking verification using general purpose proof assistants, e.g., IronFleet

Coming up with true invariants (let alone inductive ones).

- Most require an existing property, e.g., DistAI's auxiliary invariant synthesis

Getting solvers to terminate in a reasonable amount of time.

- Some argue for decision procedures that use doubly exponential time e.g., Ivy's EPR
- Some use hand-crafted heuristics and finite bounds as a best effort, e.g., FStar's fuel

Independently checking and sharing verification results.

- Most do not generate proofs (and have bugs), e.g., Dafny
- Most do not support sharing, so e.g., expensive sub-proofs cannot be offloaded to servers

Message Chains

Asynchronous Systems and Verification

Reasoning About Messages

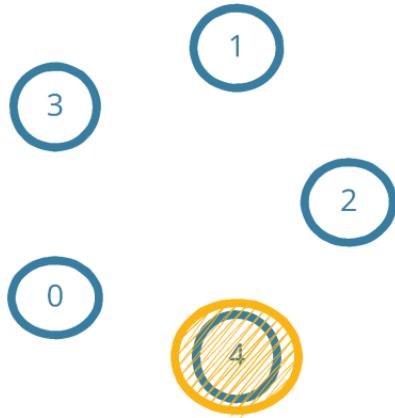
Coming Up with True Invariants

Future work: Extensions, Proofs, and Cloud Solving

Message Chains Proposal Overview

- We propose message chains
 - an organization of messages into sequences
 - that is well-aligned with existing
 - design standards and
 - programming languages
- Users can write invariants over message chains
 - making formal verification of message passing systems easier
- Later, we describe how to automatically learn some message chain invariants

Message Chains Example and Intuitive Definition



Ring Leader Election Protocol

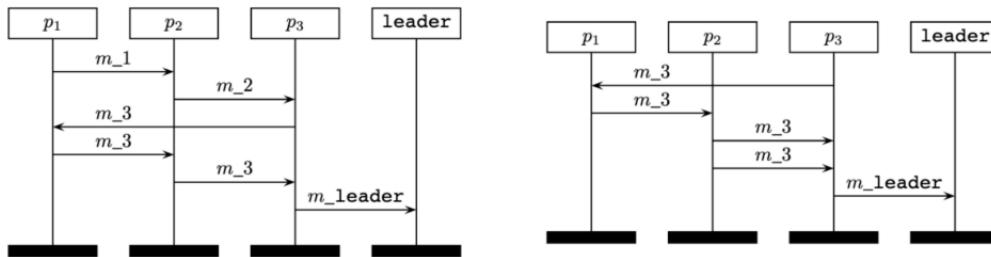
- Send your own identifier clockwise
 - starts a message chain
- If you receive a value greater than your own, send new value clockwise
 - extends a message chain
- If you receive your own identifier, declare yourself the winner

Unique Message Chains source →^{payload} target

- $0 \rightarrow^0 3;$
- $1 \rightarrow^1 2;$
- $2 \rightarrow^2 4;$
- $3 \rightarrow^3 1; 1 \rightarrow^3 2; 2 \rightarrow^3 4;$
- $4 \rightarrow^4 0; 0 \rightarrow^4 3; 3 \rightarrow^4 1; 1 \rightarrow^4 2; 2 \rightarrow^4 4;$

- If you receive and then send in the same action, then extend the message chain.
- If you send without first receiving, then start a new message chain.

Developers Rarely Think At The Level Of Individual Messages



i) One positive and one negative scenario for the leader election protocol

- "A Message Sequence Chart... describes the message flow between instances. One Message Sequence Chart describes a partial behaviour of a system."

Developers Implement Message Passing with Event Handlers

```
receive
  pong -> io:format("received pong~n", [])
end,
```

Erlang

```
def receive = {
  case "pong" => log.info("received pong")
  case _    => log.info("received unknown message")
}
```

Akka

```
on ePong do {
  print "received pong";
}
```

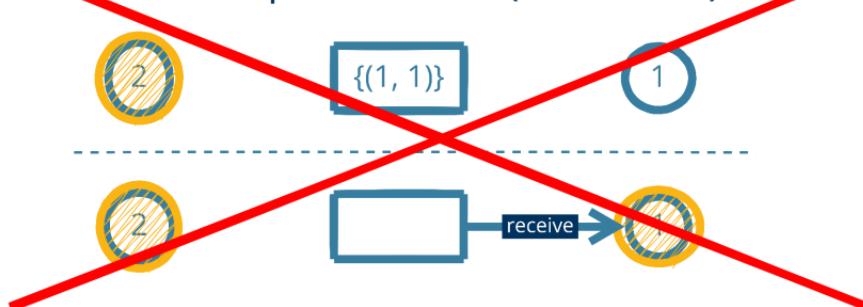
P

Ring Leader Election Protocol Verification

System Specification:

- There is always at most one node in the winning state
- For every message chain c in the current state of the system
 - the payload at the head of c is the largest source label that appears in c
 - if a node n has not yet participated in c , then n lies between the first and last nodes of c

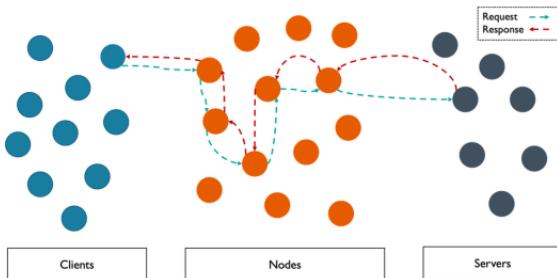
~~Old Counterexample to Induction (Now Blocked):~~



What Other Systems Does This Work For?

Example: Onion Routing Protocol

If a client receives a response, then the client must have sent a matching request.

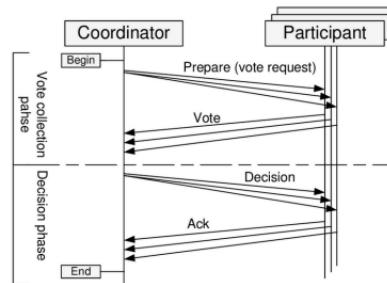


Every message chain follows the path:

Client → Node* → Server → Node* → Client

Example: Two-Phase Commit

All participants agree to either commit or abort.



Prepare → Vote → Decision → Ack

But what happens between **Vote → Decision**?

Specification Mining

Asynchronous Systems and Verification

Reasoning About Messages

Coming Up with True Invariants

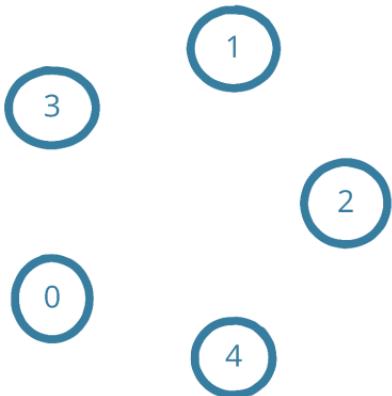
Future work: Extensions, Proofs, and Cloud Solving

What Happens When Verification Fails?



Usually the solver will timeout

Specification Mining: Can We Learn From Executions?



For every message chain c , the head payload is the largest source label that appears in c .

Algorithm

- Send your own identifier left
- If you receive a value greater than your own, send received value left
- If you receive your own identifier, declare yourself the winner

Unique Message Chains source $\rightarrow^{\text{payload}}$ target

- $0 \rightarrow^0 3;$
- $1 \rightarrow^1 2;$
- $2 \rightarrow^2 4;$
- $3 \rightarrow^3 1; 1 \rightarrow^3 2;$
- $4 \rightarrow^4 0; 0 \rightarrow^4 3; 3 \rightarrow^4 1; 1 \rightarrow^4 2; 2 \rightarrow^4 4;$

| Specification Mining, Generally

"... a machine learning approach to discovering formal specifications" that code *likely* obeys.

- Input: observed program executions
- Output: suggested invariant

The goal is not to reverse engineer the code from observations (or to discover a model).

In our setting, we are interested in mining message chain invariants:

- predicates that hold for every message chain that appears during the execution of a system

Non-Erasing Pattern Learning Example

- $\Sigma = \{C_1, C_2, C_3, S_1, S_2, S_3\}$
- $V = \{X, Y, Z\}$
- Input strings (message chains sampled from a client-server system):
 - $C_1S_1C_1$
 - $C_2S_1C_2$
 - $C_1S_2C_1$
- Angluin's output pattern: XYX
- Strings in the language of the pattern but that are kind of disappointing:
 - $S_1S_1S_1$
 - $C_2C_2S_1C_2C_2$
 - $C_1S_2C_2S_3C_3S_2C_1$

Justified Constraint Examples

- Input strings (message chains sampled from a client-server system):
 - $C_1S_1C_1$
 - $C_2S_1C_2$
 - $C_1S_2C_1$
- Output pattern: $XYX \wedge |X| = 1 \wedge |Y| = 1 \wedge X \in C^+ \wedge Y \in S^+$

Unjustified Constraint Example

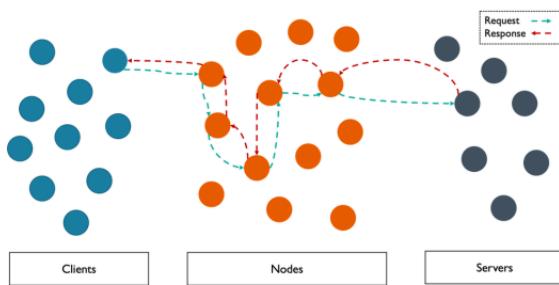
- Input strings:
 - abcd
 - dcba
 - dc
- Output pattern: XY
- Unjustified constraint (BAD): $|X| = 1$
 - e.g., the assignment $\{X \mapsto abc, Y \mapsto d\}$ can *justify* $abcd \in L(XY)$
- Note that $L(XY) = L(XY \wedge |X| = 1)$
 - We are adding *semantic* constraints to a *syntactic* notion

| Non-Erasing Patterns with Justified Constraints, Generally

- Let S be an input set of strings and p be a learned pattern
- The set of *justifications* $J(S, p)$ is the set of assignments μ such that $\exists s \in S \mu(p) = s$
- A semantic constraint is *justified* if it does not block any justification
 - In other words, a constraint f is justified iff $\forall \mu \in J(S, p) \mu \models f$
- We consider two domain-specific constraint classes
 - Single machines: $|x| = 1$, where $x \in V$
 - Machine kinds: $x \in \Gamma^+$, where $\Gamma \subseteq \Sigma$

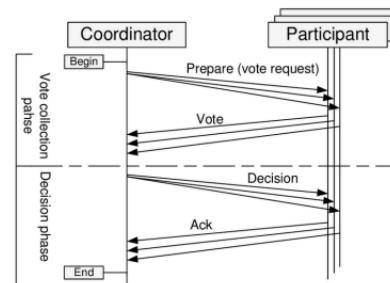
What Can Path-Patterns Learn?

Example: Onion Routing Protocol



Client Node⁺ Server Node⁺ Client
where the first and last client are the same.

Example: Two-Phase Commit



Prepare → Vote → Decision → Ack
where the three coordinator nodes are the same, but the participants before and after the decision may be different.

Learning Even More With Negative Examples

- Positive message chain examples are easy to get
 - Sample by fuzzing the target system
- Negative examples are harder
 - We *should* prove that a message chain can never happen
- PBE tools can synthesize predicates from positive and negative examples
- Assuming users can provide negative examples, PBE tools can suggest invariants
- Smyth and Burst, with less than 10 examples, can learn (among other things)
 - For every message chain c , the head payload is the largest source label that appears in c

Limitations and Future Work

Asynchronous Systems and Verification

Reasoning About Messages

Coming Up with True Invariants

Extensions, Proofs, and Cloud Solving

Challenges Revisited

- Reasoning about messages
 - Message chains help but
 - they do not remove the limitations of asynchrony
 - they do not yet support modular reasoning
 - Coming up with true invariants
 - Path pattern learning helps but captures only one class of invariants
 - Programming-by-example helps but requires negative examples
 - Independently checking verification results
 - Getting solvers to terminate
- Extensions**
- Proofs**
- Cloud Solving**

Extensions: Partial Synchrony

- Let Δ be the max message delivery time
- Let Φ be the max difference in node computation speeds
- In the asynchronous setting, no Δ or Φ exist
 - our setting so far
 - "there is no consensus protocol resilient to even one fail-stop fault"
- In the synchronous setting, Δ and Φ exist and are known
- In the partially synchronous setting, Δ and Φ exist but are not known
- What kinds of proofs can message chains help with in the partially synchronous setting?

Extensions: Modular Verification

- Decompose complex systems and reason about them in isolation
- Usually through assume-guarantee reasoning
- P supports modular testing
- Can message chains play a role in modular verification?
 - e.g., can we define sub-chains and chains of sub-chains?
 - would that help simplify proofs or improve specification mining?

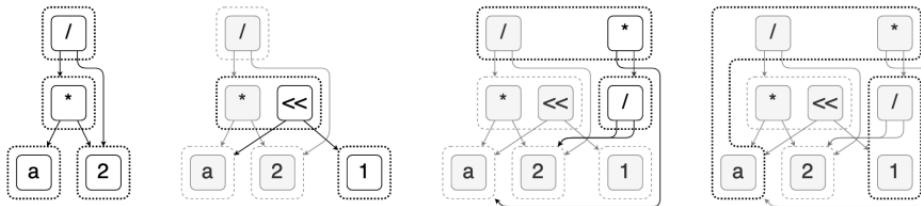
Proofs: Independently Checking Certificates

Table 1: Summary of bugs found.

	Verifier		Compiler	Other	Total
	Snd.	Prec.			
Verifier Testing	0	8	1	1	10
Compiler Testing	2	0	18	0	20
Other	1	0	0	0	1
Total	3	8	19	1	31
Confirmed	3	8	19	1	31
Fixed	0	0	7	1	8

- Modern, state-of-the-art automated verification tools, like Dafny, have bugs.
- So do SMT solvers, which are used under the hood.
- Users should not "have to rely on the correctness of a large and frequently-changing code base."

Proofs: E-Graphs to the Rescue



- E-graphs are a data structure for representing and rewriting terms
- You can extract proofs from e-graphs
- Usually input and output terms are supposed to be in the same language
 - e.g., when optimizing programs
- But can we use e-graphs as a proof generating compiler?
 - Yes! At least for some pairs of languages (which ones?)
 - But you cannot just extract the smallest term anymore

Cloud Solving: Solvers and Selection

- Different SMT solvers have different performance profiles
 - No one solver dominates all others
- Especially apparent with the emergence of cloud solvers
 - Dispatching to the cloud may take longer than solving locally
 - But may be worth it!
- Can we use e-graphs to implement proof-producing, efficient, cloud solvers?
- Can we use solver selection techniques to decide when to go to the cloud (and what solver to use)?

Learning-Enabled Verification with End-to-End Proofs

- Modular verification
 - verifying one distributed system
 - decomposed into multiple, simpler queries
- Where some queries are solved in the cloud
 - automatically selecting solver and platform per query
- With end-to-end proofs
 - verification encodings produce proofs
 - solvers produce proofs
 - proofs are stitched together and returned

Table of Desired Features, Revisited

	Domain-Specific	Implementation Level	Modular Verification	High Automation	Proof Generation	Specification Mining
TLA+	✓	✗	✓	✗	✗	✗
Ivy	✓	✗	✓	✓	✗	✗
mypyvvy	✓	✗	✗	✓	✗	✗
Coq/Lean/...	✗	✓	✓	✗	✓	✗
Dafny/F*/...	✗	✓	✓	✓	✗	✗
Proposal	✓	✓	✓*	✓	✓*	✓

Timeline

- Message Chains for Asynchronous Verification
 - April 2023
 - submit to SOSP
- Learning Invariants using Message Chain Examples
 - May 2023
 - submit to FMCAD
- Modular Message Chains and Partially Synchronous Systems
 - September 2023 to December 2023
 - formalize modular message chains
 - encode partial synchrony in our framework
 - January 2024 to March 2024
 - challenge: verify one system per week of 294-234
 - ~April 2024
 - submit to SOSP or similar
- SMT Solver and Platform Selection: When and How to Cloud
 - March 2023 to September 2023
 - proof-of-concept implementation
 - e-graph-based eager SMT solving in the cloud
 - algorithm selection for when to go to the cloud
 - January 2024
 - submit to CAV
- Proof-Producing Verification Encodings with Learned Rewrites
 - ~May 2024 to December 2024
 - combine previous projects
 - January 2025
 - submit to CAV
- Write thesis: January 2025 to May 2025
- Graduation: May 2025